WorkCoPilot - Novel Patent Claims

COE Security LLC / AdmireVeda | 2026-10-25 | 14 novel elements, mapped to the 20-claim set (3 independent: 1 / 10 / 16)

DRAFT - verify with a registered practitioner; prior-art search is non-exhaustive. Reformatted from patent-package/02 and /04; no new prior art, claims or facts added.

No.	Novel element	What it is (plain language)	Why it is novel (vs the prior art found)	Backed by (independent + dependent claim numbers)
1	**Read-only guidance - the system never injects input**	The client only displays the next instruction; it never sends a click, keystroke, command, or form submission into the application you are using. You perform every action yourself, so the human stays accountable for every change.	Agentic / RPA assistants (Verizon US 2025/0238087, Apple US 12,406,664, Vasylyev US 2024/0412720) act for the user, and even the closest guided-help art (USAA US 11,989,786 and US 12,346,977) recites a system that "automatically assists." No verified reference claims a screen-grounded assistant structurally barred from injecting input, with the human's own action as the only state change.	Independent: 1, 10, 16. Dependent: 20.
2	**Observe -> instruct -> verify -> correct -> advance loop grounded in real-time screen state**	The plan moves to the next step only after a later captured frame actually shows the expected result; a wrong move produces a corrective instruction instead of advancing.	Prior guided-help (Boomersurf US 10,852,898; USAA US 11,989,786 / US 12,346,977; Meta US 11,036,984; Visautact US 10,290,222) gates on GUI-object/accessibility data, camera/AR input, or pre-authored tutorials - not vision over a screenshot of an arbitrary application used as the verification oracle.	Independent: 1, 10, 16. Dependent: 2, 6.
3	**Visible-when-done condition on every plan step**	The model-generated checklist attaches to each step a description of exactly what will be on screen when that step is truly finished, so completion is read from the display rather than assumed.	Prior art relies on pre-authored step images or sensed micro-steps. No verified reference pairs a model-generated plan for an arbitrary user-stated goal with a per-step visible-when-done test evaluated from a redacted screenshot.	Independent: 1, 10, 16.
4	**Structured per-frame navigation record**	For each frame the vision model returns a single record: completed steps, the one current step, one screen-specific micro-action, a one-line observation, a problem flag, and a done flag.	No verified reference defines this specific completed / current / guidance / observation / flag / done contract, exchanged once per frame, as the control structure for guidance the human then performs.	Independent: 1, 10, 16. Dependent: 2.
5	**Proactive prerequisite insertion and UI-drift re-targeting**	If something the current step needs is missing on screen, the guidance switches to establishing it first; if a named control has moved or been renamed since planning, the guidance points to the equivalent visible control and states that the interface differs.	Pre-authored walkthroughs (Boomersurf; WalkMe-class digital adoption platforms) break when the UI drifts or a prerequisite is absent; they cannot re-target to an equivalent control seen in a live frame or insert a missing prerequisite on the fly.	Dependent: 6.
6	**On-device OCR frame redaction before any frame is transmitted**	Before a frame leaves the device, on-device OCR reads each text line and the same policy engine used for typed text masks every line holding sensitive data; the masked image is what is sent.	Capital One US 12,013,965 masks only its own application's known GUI fields; JPMorgan US 2025/0005060, IBM US 2025/0086310 and others redact at a server/gateway. None OCRs an arbitrary third-party app's frame on-device using the identical DLP engine (including tenant block-words and custom patterns) as the text path.	Independent: 1, 10, 16. Dependent: 3, 14.
7	**Fail-safe downscale backstop when OCR is unavailable**	If the on-device OCR engine cannot run, the frame is heavily downscaled so fine text is illegible rather than transmitted at full resolution - a safe default, honestly reported as a best-effort backstop.	No verified redaction reference describes a legibility-destroying downscale fallback that prevents full-resolution leakage when the on-device recognizer fails to load or run.	Dependent: 4.
8	**Reversible, device-only tokenization with a volatile token map**	Sensitive values in outbound text are replaced with stable typed placeholders (for example [[EMAIL_1]], [[IP_2]]); the token-to-value map lives only in device memory, is never transmitted or written to any server log, and is discarded when the session ends.	Capital One US 2024/0111891 sanitizes irreversibly; JPMorgan US 2025/0005060 and IBM US 2025/0086310 mask at a server / moderator with no restoration. No verified reference keeps a reversible mapping device-only for later local restoration in a screen-guidance client.	Independent: 10. Dependent: 5, 11.

No.	Novel element	What it is (plain language)	Why it is novel (vs the prior art found)	Backed by (independent + dependent claim numbers)
9	**Local rehydration of the model's answer, including streaming**	Placeholders in the model's reply are swapped back to the real values on the device - applied to each streamed chunk as it arrives - so the user reads a complete answer that the server and the model never saw.	Server- and moderator-side masking art does not restore the originals in the response at all. No verified reference performs on-device, streaming rehydration of a model answer whose sensitive values never left the device.	Independent: 10. Dependent: 5, 12.
10	**Never-guess fallback for a missing or garbled token**	If a placeholder is missing or altered in the reply, the text is left as the placeholder and flagged incomplete at that point - the system never invents or substitutes a value not in the token map.	Reversible-masking schemes do not describe an explicit no-fabrication rule that leaves and flags the placeholder rather than guessing a value, protecting against silent corruption of the restored answer.	Dependent: 13.
11	**Deterministic destructive-command and vulnerable-dependency gate**	Every command or code block in a plan is matched against a catalog of destructive patterns and flagged (hidden behind an explicit acknowledgment), and named dependency versions are checked against a vulnerability database before the user is shown them.	Guided-help and agentic art do not add a deterministic safety gate over the exact commands and code handed to a human, with a report-rather-than-silently-pass behaviour when the vulnerability source is unreachable.	Dependent: 15.
12	**Evidence captured only on verification, reusing the already-redacted frame**	When a step is first verified, the frame already redacted for transmission is kept as that step's evidence - once per step, encrypted at rest - so nothing additional ever leaves the device to build the record.	Screen-recording / DLP art captures a separate recording. No verified reference ties evidence capture to the verification event and reuses the same already-redacted frame, adding zero additional egress.	Independent: 16. Dependent: 8, 9.
13	**Tenant-enforced egress policy including the screenshot-evidence retention switch**	A per-company policy object (per-category send / redact / withhold / tokenize, block-words, mode, a switch for whether redacted evidence frames are kept at all, a retention window, armed/locked) is pushed to member devices and enforced there; members cannot weaken it and the server re-applies it as defense in depth.	General cloud / endpoint DLP exists, but no verified reference binds a per-tenant policy object - including an on/off switch for retaining screenshot evidence - to on-device enforcement inside an AI screen-guidance client with arm and lock semantics a member cannot override.	Independent: 16. Dependent: 17, 18.
14	**Audit log recording each redacted transmission by counts, not the removed values**	Every transmission is logged with the policy version in force, the mode, counts of what was removed per category, and whether a frame was sent - never the removed values; a sweep then deletes each expired session together with its evidence frames and its audit log in one operation.	No verified reference records a per-transmission audit that deliberately stores counts rather than the sensitive values, coupled to a single-operation retention sweep of the session, its evidence, and its log together.	Independent: 16. Dependent: 19.